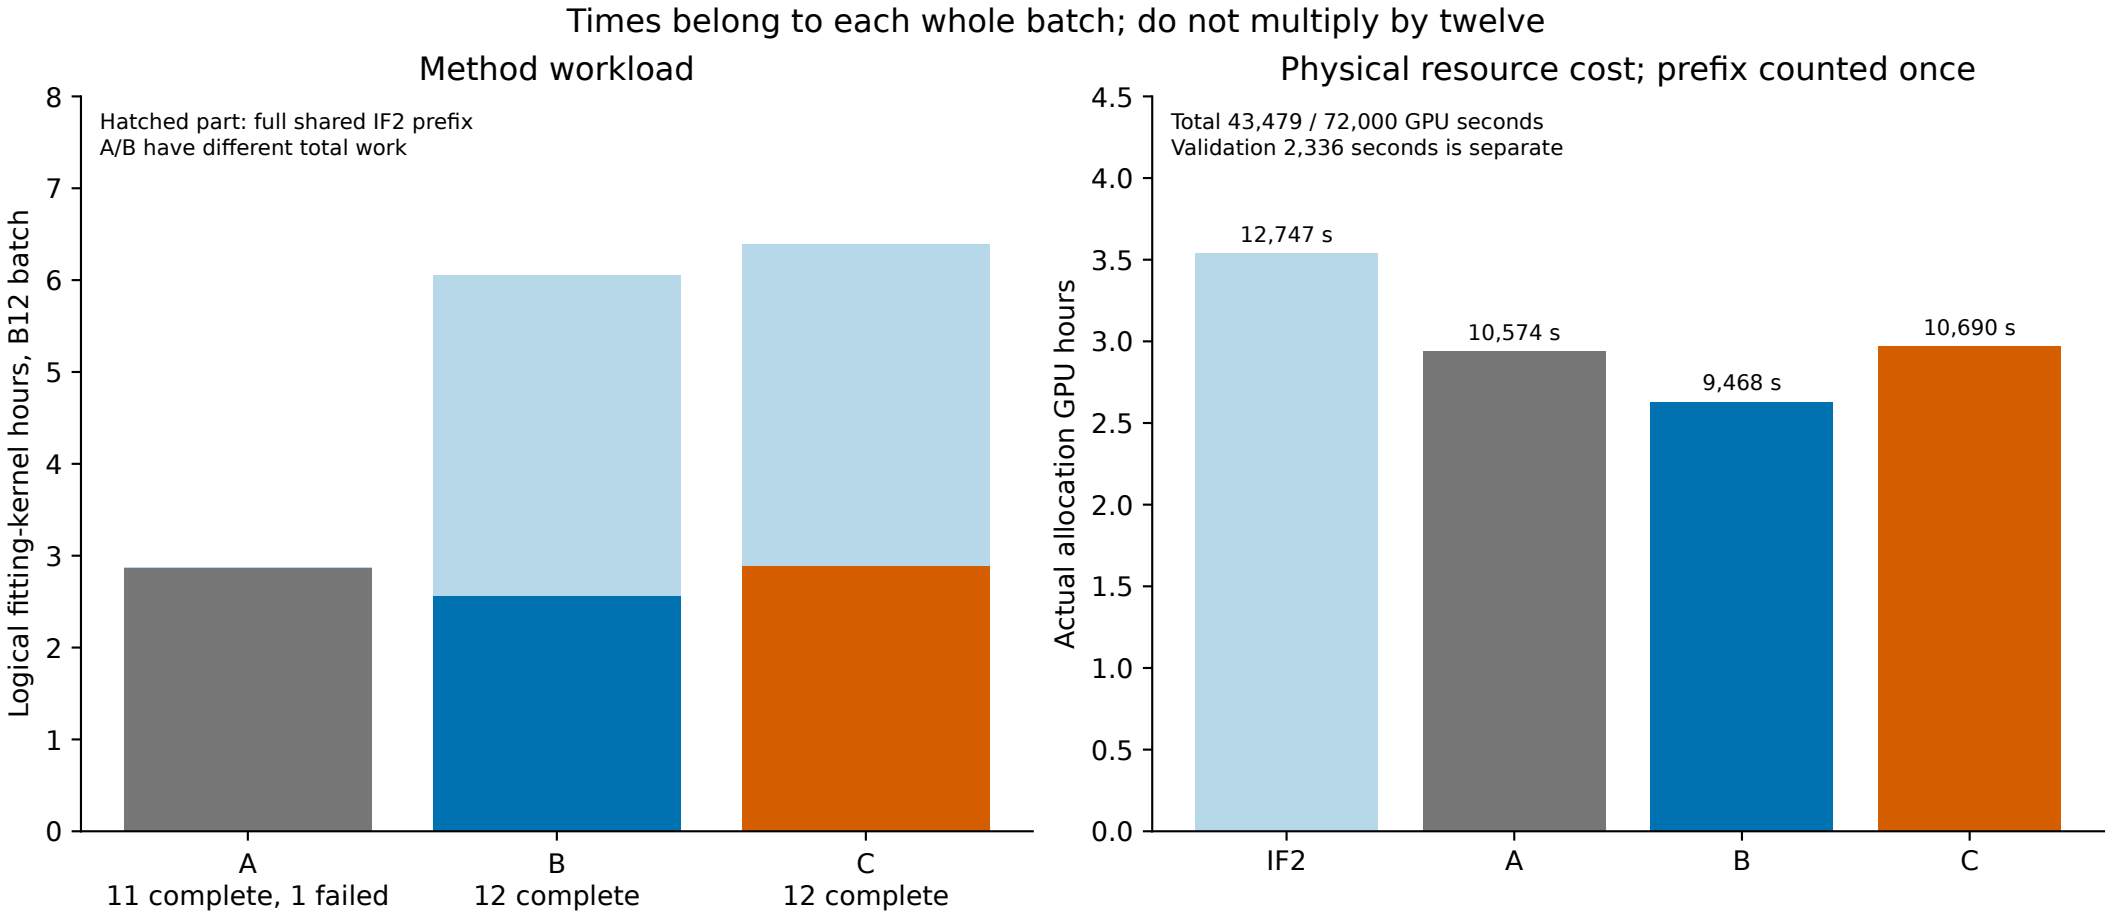
Times belong to each whole batch; do not multiply by twelve
Method workload
Hatched part: full shared IF2 prefix
A/B have different total work
Logical fitting-kernel hours, B12 batch
0
1
2
3
4
5
6
7
8
A
11 complete, 1 failed
B
12 complete
C
12 complete
Physical resource cost; prefix counted once
Total 43,479 / 72,000 GPU seconds
Validation 2,336 seconds is separate
Actual allocation GPU hours
0.0
0.5
1.0
1.5
2.0
2.5
3.0
3.5
4.0
4.5
12,747 s
10,574 s
9,468 s
10,690 s
IF2
A
B
C